

# ChurnPrediction

October 13, 2023

## 0.1 Introduction

Subscription services are leveraged by companies across many industries, from fitness to video streaming to retail. One of the primary objectives of companies with subscription services is to decrease churn and ensure that users are retained as subscribers. In order to do this efficiently and systematically, many companies employ machine learning to predict which users are at the highest risk of churn, so that proper interventions can be effectively deployed to the right audience.

In this project, I will be tackling the churn prediction problem on a very unique and interesting group of subscribers on a video streaming service.

### 0.1.1 Train vs. Test

The data is available [here](#).

The data is already split into train and test set.

`train.csv` contains 70% of the overall sample (243,787 subscriptions to be exact) and importantly, will reveal whether or not the subscription was continued into the next month (the “ground truth”).

The `test.csv` dataset contains the exact same information about the remaining segment of the overall sample (104,480 subscriptions to be exact) except the actual churn data. The goal of this project is to make prediction on each customer’s churn.

### 0.1.2 Dataset descriptions

```
[1]: import pandas as pd
data_descriptions = pd.read_csv('data_descriptions.csv')
pd.set_option('display.max_colwidth', None)
data_descriptions
```

```
[1]:
```

	Column_name	Column_type	Data_type	\
0	AccountAge	Feature	integer	
1	MonthlyCharges	Feature	float	
2	TotalCharges	Feature	float	
3	SubscriptionType	Feature	object	
4	PaymentMethod	Feature	string	
5	PaperlessBilling	Feature	string	

6	ContentType	Feature	string
7	MultiDeviceAccess	Feature	string
8	DeviceRegistered	Feature	string
9	ViewingHoursPerWeek	Feature	float
10	AverageViewingDuration	Feature	float
11	ContentDownloadsPerMonth	Feature	integer
12	GenrePreference	Feature	string
13	UserRating	Feature	float
14	SupportTicketsPerMonth	Feature	integer
15	Gender	Feature	string
16	WatchlistSize	Feature	float
17	ParentalControl	Feature	string
18	SubtitlesEnabled	Feature	string
19	CustomerID	Identifier	string
20	Churn	Target	integer

	Description	
0		The age of the user's account in months.
1		The amount charged to the user on a monthly basis.
2		The total charges incurred by the user over the account's lifetime.
3		The type of subscription chosen by the user (Basic, Standard, or Premium).
4		The method of payment used by the user.
5		Indicates whether the user has opted for paperless billing (Yes or No).
6		The type of content preferred by the user (Movies, TV Shows, or Both).
7		Indicates whether the user has access to the service on multiple devices (Yes or No).
8		The type of device registered by the user (TV, Mobile, Tablet, or Computer).
9		The number of hours the user spends watching content per week.
10		The average duration of each viewing session in minutes.
11		The number of content downloads by the user per month.
12		The preferred genre of content chosen by the user.
13		The user's rating for the service on a scale of 1 to 5.
14		The number of support tickets raised by the user per month.

15		The gender of
the user (Male or Female).		
16		The number of items
in the user's watchlist.		
17	Indicates whether parental control is enabled	
for the user (Yes or No).		
18	Indicates whether subtitles are enabled	
for the user (Yes or No).		
19		A unique
identifier for each customer.		
20	The target variable indicating whether a user has churned or not (1 for	
churned, 0 for not churned).		

## 0.2 Setup

```
[3]: import pandas as pd
import numpy as np
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.dummy import DummyClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, \
    recall_score, ConfusionMatrixDisplay
from scipy.stats import randint
from matplotlib import pyplot as plt
import seaborn as sns
import plotly.express as px
%matplotlib inline
```

## 0.3 Load the Data

```
[4]: train_df = pd.read_csv("train.csv")
print('train_df Shape:', train_df.shape)
train_df.head()
```

train\_df Shape: (243787, 21)

```
[4]: AccountAge MonthlyCharges TotalCharges SubscriptionType \
0      20      11.055215      221.104302      Premium
1      57       5.175208      294.986882       Basic
2      73      12.106657      883.785952       Basic
3      32       7.263743      232.439774       Basic
4      57      16.953078      966.325422      Premium

PaymentMethod PaperlessBilling ContentType MultiDeviceAccess \
```

0	Mailed check	No	Both	No
1	Credit card	Yes	Movies	No
2	Mailed check	Yes	Movies	No
3	Electronic check	No	TV Shows	No
4	Electronic check	Yes	TV Shows	No

	DeviceRegistered	ViewingHoursPerWeek	...	ContentDownloadsPerMonth	\
0	Mobile	36.758104	...		10
1	Tablet	32.450568	...		18
2	Computer	7.395160	...		23
3	Tablet	27.960389	...		30
4	TV	20.083397	...		20

	GenrePreference	UserRating	SupportTicketsPerMonth	Gender	WatchlistSize	\
0	Sci-Fi	2.176498		4	Male	3
1	Action	3.478632		8	Male	23
2	Fantasy	4.238824		6	Male	1
3	Drama	4.276013		2	Male	24
4	Comedy	3.616170		4	Female	0

	ParentalControl	SubtitlesEnabled	CustomerID	Churn
0	No	No	CB6SXPNVZA	0
1	No	Yes	S7R2G87009	0
2	Yes	Yes	EASDC20BDT	0
3	Yes	Yes	NPF69NT69N	0
4	No	No	4LGYPK7VOL	0

[5 rows x 21 columns]

```
[5]: test_df = pd.read_csv("test.csv")
      print('test_df Shape:', test_df.shape)
```

test\_df Shape: (104480, 20)

## 0.4 Preprocess & EDA

```
[6]: train_df.isna().any()
```

```
[6]: AccountAge           False
      MonthlyCharges       False
      TotalCharges         False
      SubscriptionType     False
      PaymentMethod        False
      PaperlessBilling      False
      ContentType          False
      MultiDeviceAccess     False
```

```

DeviceRegistered      False
ViewingHoursPerWeek   False
AverageViewingDuration False
ContentDownloadsPerMonth False
GenrePreference       False
UserRating            False
SupportTicketsPerMonth False
Gender               False
WatchlistSize         False
ParentalControl       False
SubtitlesEnabled      False
CustomerID            False
Churn                 False
dtype: bool

```

```
[7]: train_df.describe()
```

```

[7]:      AccountAge  MonthlyCharges  TotalCharges  ViewingHoursPerWeek  \
count  243787.000000  243787.000000  243787.000000  243787.000000
mean    60.083758    12.490695    750.741017    20.502179
std     34.285143     4.327615    523.073273    11.243753
min      1.000000     4.990062     4.991154     1.000065
25%     30.000000     8.738543    329.147027    10.763953
50%     60.000000    12.495555    649.878487    20.523116
75%     90.000000    16.238160   1089.317362    30.219396
max    119.000000    19.989957   2378.723844    39.999723

      AverageViewingDuration  ContentDownloadsPerMonth  UserRating  \
count      243787.000000      243787.000000  243787.000000
mean         92.264061         24.503513      3.002713
std         50.505243         14.421174      1.155259
min          5.000547          0.000000      1.000007
25%         48.382395         12.000000      2.000853
50%         92.249992         24.000000      3.002261
75%        135.908048         37.000000      4.002157
max        179.999275         49.000000      4.999989

      SupportTicketsPerMonth  WatchlistSize  Churn
count      243787.000000  243787.000000  243787.000000
mean         4.504186     12.018508     0.181232
std         2.872548      7.193034     0.385211
min          0.000000      0.000000      0.000000
25%          2.000000      6.000000      0.000000
50%          4.000000     12.000000      0.000000
75%          7.000000     18.000000      0.000000
max          9.000000     24.000000      1.000000

```

```
[8]: train_df.dtypes
```

```
[8]: AccountAge           int64
MonthlyCharges         float64
TotalCharges           float64
SubscriptionType       object
PaymentMethod          object
PaperlessBilling       object
ContentType            object
MultiDeviceAccess      object
DeviceRegistered       object
ViewingHoursPerWeek    float64
AverageViewingDuration float64
ContentDownloadsPerMonth int64
GenrePreference        object
UserRating             float64
SupportTicketsPerMonth int64
Gender                 object
WatchlistSize          int64
ParentalControl        object
SubtitlesEnabled       object
CustomerID             object
Churn                  int64
dtype: object
```

Check the classes for each categorical variable.

```
[9]: train_df["SubscriptionType"].unique()
```

```
[9]: array(['Premium', 'Basic', 'Standard'], dtype=object)
```

```
[10]: train_df["PaymentMethod"].unique()
```

```
[10]: array(['Mailed check', 'Credit card', 'Electronic check', 'Bank transfer'],
        dtype=object)
```

```
[11]: train_df["PaperlessBilling"].unique()
```

```
[11]: array(['No', 'Yes'], dtype=object)
```

```
[12]: train_df["ContentType"].unique()
```

```
[12]: array(['Both', 'Movies', 'TV Shows'], dtype=object)
```

```
[13]: train_df["MultiDeviceAccess"].unique()
```

```
[13]: array(['No', 'Yes'], dtype=object)
```

```
[14]: train_df["DeviceRegistered"].unique()
```

```
[14]: array(['Mobile', 'Tablet', 'Computer', 'TV'], dtype=object)
```

```
[15]: train_df["GenrePreference"].unique()
```

```
[15]: array(['Sci-Fi', 'Action', 'Fantasy', 'Drama', 'Comedy'], dtype=object)
```

```
[16]: train_df["Gender"].unique()
```

```
[16]: array(['Male', 'Female'], dtype=object)
```

```
[17]: train_df["ParentalControl"].unique()
```

```
[17]: array(['No', 'Yes'], dtype=object)
```

```
[18]: train_df["SubtitlesEnabled"].unique()
```

```
[18]: array(['No', 'Yes'], dtype=object)
```

Now, dummy code the categorical variables except CustomerID.

```
[19]: train_df_new = train_df
train_df_new = train_df.drop("CustomerID", axis=1)
```

```
[20]: categorical_columns = ['SubscriptionType', 'PaymentMethod', 'PaperlessBilling',
    ↳ 'ContentType', 'MultiDeviceAccess',
    ↳ 'DeviceRegistered', 'GenrePreference', 'Gender',
    ↳ 'ParentalControl', 'SubtitlesEnabled']

for col in categorical_columns:
    dummies = pd.get_dummies(train_df_new[col], prefix=col)
    train_df_new = pd.concat([train_df_new, dummies], axis=1)

train_df_new = train_df_new.drop(categorical_columns, axis=1)
```

Check the proportion of churn in the training set.

```
[21]: fig = px.histogram(train_df['Churn'].astype(str), color=train_df['Churn'].
    ↳ astype(str), title="Churn count")

# Calculate percentages
total_count = len(train_df)
percentage_labels = [(f'{churn}: {count / total_count:.2%}')] for churn, count_
    ↳ in train_df['Churn'].value_counts().items()

# Add percentage as a caption
```

```
fig.update_layout(
    annotations=[
        dict(
            text=', '.join(percentage_labels),
            showarrow=False,
            x=0.5,
            y=1.1,
            font=dict(size=12),
        )
    ]
)

# Show the plot
fig.show()
```

## 0.5 Make predictions (required)

In this project, I decided to use random forest classifier because \* the dataset is pretty large \* the dataset has many variables to be considered \* I just learned in my class and wanted to try

```
[22]: X_train = train_df_new.drop("Churn", axis = 1)
      y_train = train_df_new["Churn"]
      model = RandomForestClassifier()
      model.fit(X_train, y_train)
```

```
[22]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                             criterion='gini', max_depth=None, max_features='auto',
                             max_leaf_nodes=None, max_samples=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=100,
                             n_jobs=None, oob_score=False, random_state=None,
                             verbose=0, warm_start=False)
```

```
[23]: feature_importances = pd.Series(model.feature_importances_, index=X_train.
      ↪columns).sort_values(ascending=False)


      fig = px.bar(feature_importances)
      fig.show()
```

There are clear differences between the important and non important variables. For the ones with importance > 0.04, check if there are any correlations between one another.

```
[24]: plt.figure(figsize=(16, 6))
      important_variables = feature_importances[0:9].index
      sns.heatmap(train_df[important_variables].corr())
```



```
[24]: <matplotlib.axes._subplots.AxesSubplot at 0x7f76f90831d0>
```



The heatmap shows that there are correlations between “AccountAge” and “TotalCharges,” and “MonthlyCharges” and “TotalCharges” which makes a lot of sense. Remove all the unimportant variables as well as “TotalCharges” from the model and run a new model.

```
[25]: important_variables = important_variables.drop("TotalCharges")
X_train = train_df_new[important_variables]
y_train = train_df_new["Churn"]
model2 = RandomForestClassifier()
model2.fit(X_train, y_train)
```

```
[25]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                             criterion='gini', max_depth=None, max_features='auto',
                             max_leaf_nodes=None, max_samples=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=100,
                             n_jobs=None, oob_score=False, random_state=None,
                             verbose=0, warm_start=False)
```

```
[26]: feature_importances = pd.Series(model2.feature_importances_, index=X_train.
    ↪columns).sort_values(ascending=False)

fig = px.bar(feature_importances)
fig.show()
```

Now that we identify the variables to include, we will move on to hypertune the model by determining the best parameters

```
[27]: # param_dist = {'n_estimators': randint(50,500),  
#               'max_depth': randint(1,20)}
```

```
# rf = RandomForestClassifier()
```

```
# rand_search = RandomizedSearchCV(rf,  
#                               param_distributions = param_dist,  
#                               n_iter=5,  
#                               cv=5)
```

```
# rand_search.fit(X_train, y_train)
```

```
[28]: # best_rf = rand_search.best_estimator_
```

```
# print('Best hyperparameters:', rand_search.best_params_)
```

```
[29]: best_rf = RandomForestClassifier(max_depth = 13, n_estimators = 285)  
best_rf.fit(X_train, y_train)
```

```
[29]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,  
                             criterion='gini', max_depth=13, max_features='auto',  
                             max_leaf_nodes=None, max_samples=None,  
                             min_impurity_decrease=0.0, min_impurity_split=None,  
                             min_samples_leaf=1, min_samples_split=2,  
                             min_weight_fraction_leaf=0.0, n_estimators=285,  
                             n_jobs=None, oob_score=False, random_state=None,  
                             verbose=0, warm_start=False)
```

Using the model with the best parameters, we will finally make predictions on the test data.

```
[30]: X_test = test_df[important_variables]  
y_pred = best_rf.predict(X_test)  
prediction_df = pd.DataFrame({'CustomerID': test_df['CustomerID'], 'preds':  
    ↪ y_pred})
```

## 0.6 Post-process

```
[31]: fig = px.histogram(prediction_df.preds)  
fig.show()
```

```
[32]: prediction_df.preds.value_counts()
```

```
[32]: 0    102111  
     1     2369  
     Name: preds, dtype: int64
```

```
[33]: result = test_df
      result['preds'] = prediction_df['preds']
```

Check how predicted churn is different for the variables with importance

```
[34]: fig = px.box(x = result.preds, y = result.AverageViewingDuration)
      fig.show()
```

You can see that the users who churned spend shorter amount of time watching a show/movie. Possible explanations are \* they do not like the contents on this platform or do not have time to watch, resulting in giving up before finishing them \* they prefer the contents of shorter duration (short films, an episode of a show)

```
[35]: fig = px.box(x = result.preds, y = result.ViewingHoursPerWeek)
      fig.show()
```

Similarly, the churned users spent significantly less time using the streaming service.

```
[36]: fig = px.box(x = result.preds, y = result.MonthlyCharges)
      fig.show()
```

The churn group pays higher subscription fee, which is a reasonable incentive to cancel the plan.

```
[37]: fig = px.box(x = result.preds, y = result.AccountAge)
      fig.show()
```

The users in the churned group did not use the platform for a long period of time.

```
[38]: fig = px.box(x = result.preds, y = result.ContentDownloadsPerMonth)
      fig.show()
```

The users in the churned group do not take advantage of content downloading or might not be aware of the functionality, which explains their dissatisfaction.

```
[39]: fig = px.box(x = result.preds, y = result.SupportTicketsPerMonth)
      fig.show()
```

The number of support tickets are larger for the churned group. We can assume that they had trouble with the service or concern about their plan.

## 0.7 Conclusion & Next Step

In this project, Random Forest Classifier was used to predict the customer churn on subscription service. First, the exploratory data analysis showed the potential relationship with the variables and the churn. Then, the model was developed and fine tuned. The model used average viewing duration, viewing hours per week, monthly charges, account age, user rating, content downloads per month, watchlist size, and support tickets per month as the variables with high importance. Since the test data does not contain the correct predictions, unfortunately, we were unable to see the prediction accuracy and the model performance.

As the next step, we would like to test out different models, such as logistic regression, k-nearest neighbors, decision tree classifier, and support vector machine, and compare the model performance.